

Assignment - Classification Algorithm

Problem Statement:

Using patient clinical and laboratory measurements, we want to automatically predict whether a patient has chronic kidney disease (CKD) or not (classification: yes/no).

Problem Identification:

Machine Learning Problem

- Type: Supervised Learning – **Binary Classification**.
- Goal: Build an accurate, reliable classifier that:
 - Correctly identifies CKD patients (high recall for CKD class)
 - Keeps false alarms low (reasonable precision)
 - Has good overall performance (accuracy, F1-score)

Dataset Details:

Total Columns: 25 (Laboratory Measurements)

- Features (24):
age, bp, sg, al, su, rbc, pc, pcc, ba, bgr, bu, sc, sod, pot, hrmo, pcv, wc, rc, htn, dm, cad, appet, pe, ane
- Target (1):
Classification

Total Rows: 399 (Patient Records)

Feature types:

- **Numeric features:**

- age: age in years
- bp: blood pressure
- al, su: albumin, sugar
- bgr, bu, sc, sod, pot: various biochemical test results
- hrmo, pcv, wc, rc: blood count parameters
- **Categorical features:**
 - sg: takes values like a, b, c, d, e
 - rbc, pc: normal, abnormal
 - pcc, ba: present, notpresent
 - htn, dm, cad, pe: yes, no
 - appet, ane: good, poor
- **Target:**
 - classification: yes (CKD) / no (non-CKD)

Data Preprocessing:

Among the input variables, 11 features are **Nominal Data**. To convert this, **One-hot Encoding** technique is used.

- **Binary yes/no columns:**
 - htn, dm, cad, pe, ane, classification: yes → 1, no → 0
- **Binary normal/abnormal:**
 - rbc, pc: abnormal → 1, normal → 0
- **Binary present/notpresent:**
 - pcc, ba: present → 1, notpresent → 0
- **Binary appetite:**
 - appet: poor → 1, good → 0 (if good existed)

One feature has Ordinal Data and **Label Encoding** is used to convert this into numerical values.

- **ordinal-like feature sg (a, b, c, d, e):**
 - Simple **label encoding**: a → 0, b → 1, c → 2, d → 3, e → 4

For **distance-based / linear models** (Logistic Regression, SVM, KNN), StandardScaler is used for **Data Standardization**.

Evaluation metrics across various algorithms for the dataset:

1. Logistic Regression

```
Best Logistic Regression Params: {'penalty': 'l2', 'solver': 'newton-cg'}
```

```
===== Logistic Regression =====
```

```
Accuracy: 0.98
```

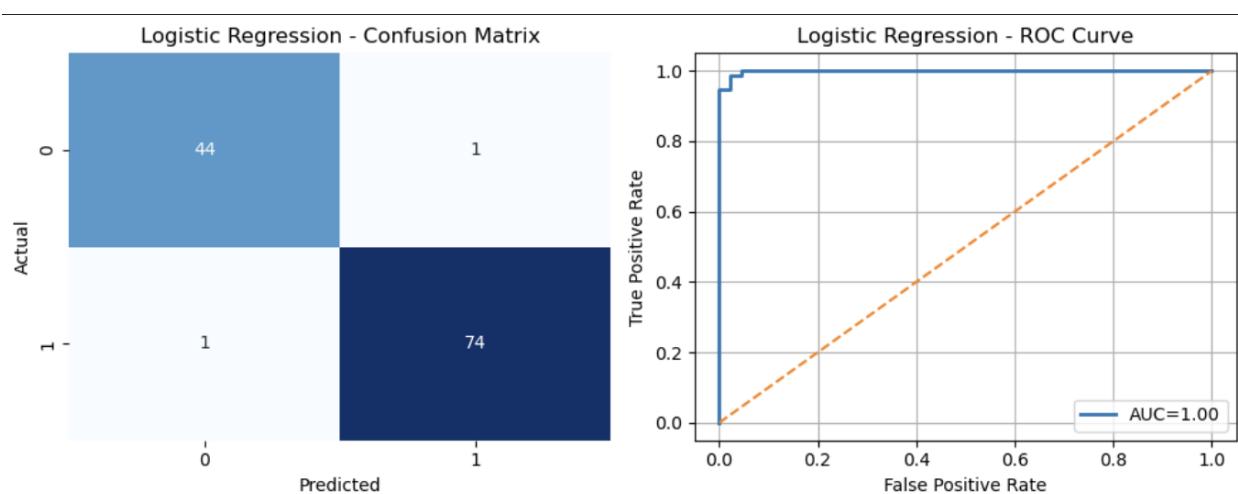
```
Confusion Matrix:
```

```
[[44  1]
 [ 1 74]]
```

```
Classification report:
```

	precision	recall	f1-score	support
0	0.98	0.98	0.98	45
1	0.99	0.99	0.99	75
accuracy			0.98	120
macro avg	0.98	0.98	0.98	120
weighted avg	0.98	0.98	0.98	120

```
AUC Score: 1.00
```

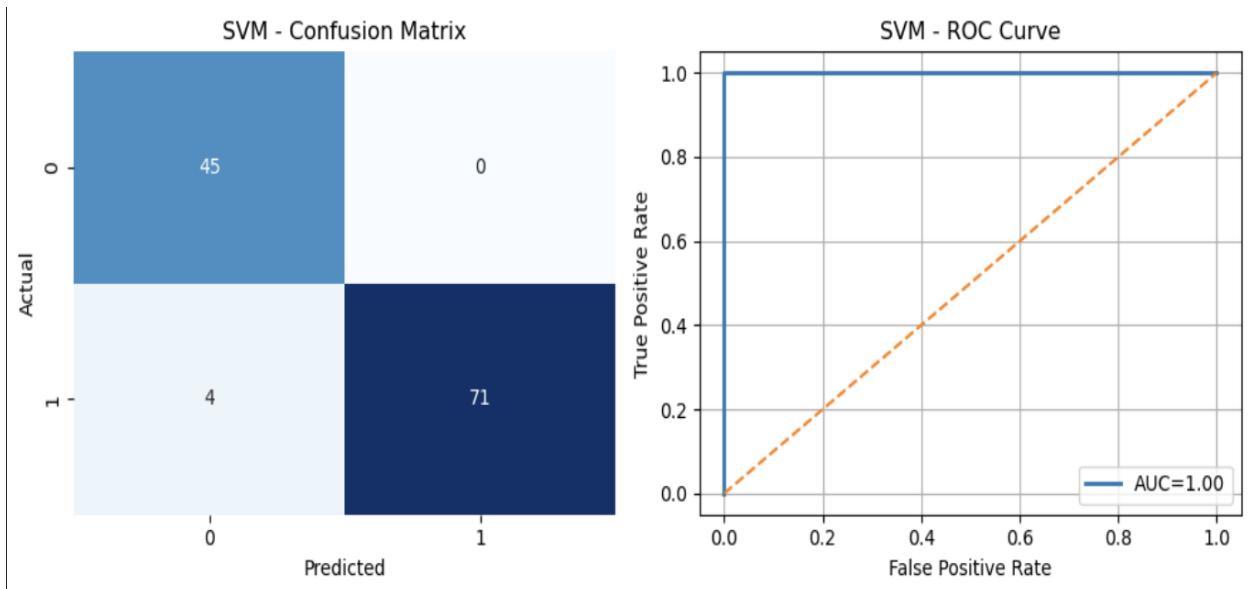


2. Support Vector Machine:

```
Best SVM Params: {'C': 0.1, 'gamma': 'scale', 'kernel': 'rbf'}
```

===== SVM =====

```
Accuracy: 0.97
Confusion Matrix:
[[45  0]
 [ 4 71]]
Classification report:
      precision    recall  f1-score   support
          0       0.92     1.00     0.96      45
          1       1.00     0.95     0.97      75
                                              accuracy         0.97      120
                                              macro avg     0.96      120
                                              weighted avg  0.97      120
AUC Score: 1.00
```



3. Decision Tree:

```
Best Decision Tree Params: {'criterion': 'entropy', 'max_features': 'sqrt', 'splitter': 'random'}
```

```
===== Decision Tree =====
```

```
Accuracy: 0.96
```

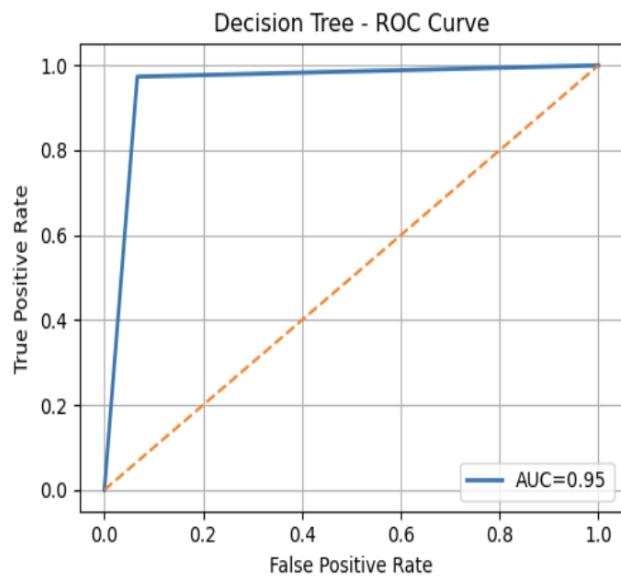
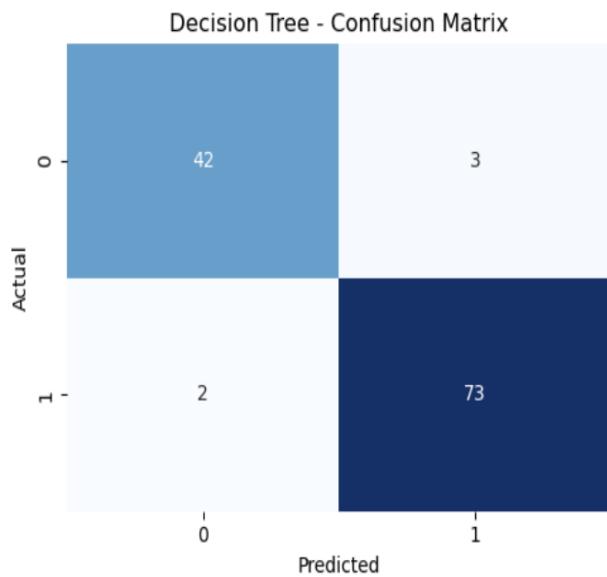
```
Confusion Matrix:
```

```
[[42  3]
 [ 2 73]]
```

```
Classification report:
```

	precision	recall	f1-score	support
0	0.95	0.93	0.94	45
1	0.96	0.97	0.97	75
accuracy			0.96	120
macro avg	0.96	0.95	0.96	120
weighted avg	0.96	0.96	0.96	120

```
AUC Score: 0.95
```



4. Random Forest:

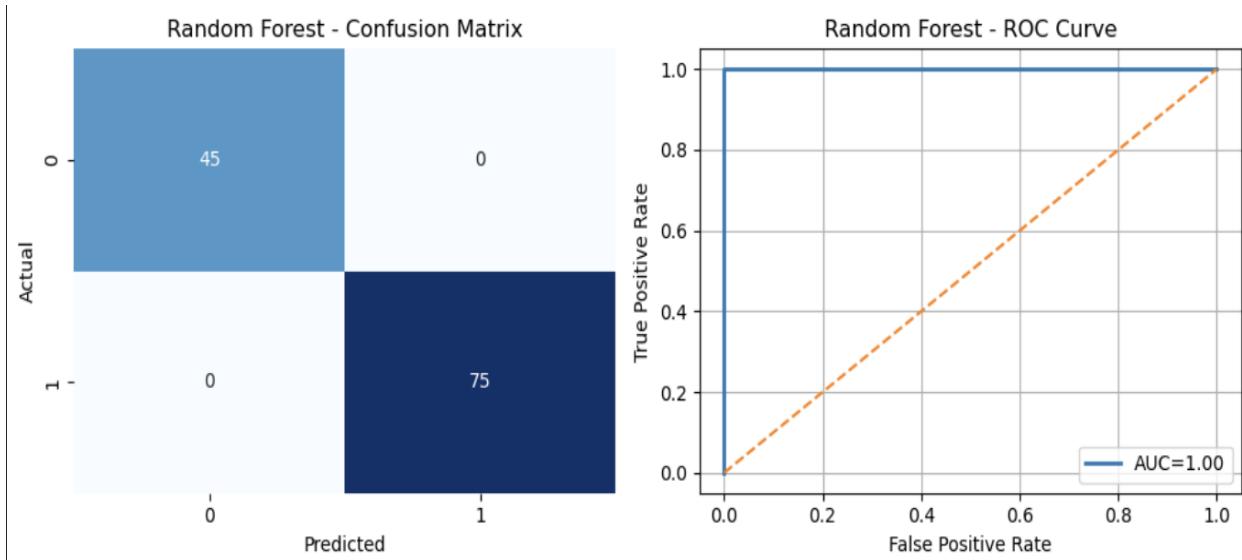
```
Best Random Forest Params: {'criterion': 'gini', 'max_features': 'sqrt', 'n_estimators': 100}

===== Random Forest =====
Accuracy: 1.00
Confusion Matrix:
[[45  0]
 [ 0 75]]
Classification report:
      precision    recall  f1-score   support

          0       1.00     1.00     1.00      45
          1       1.00     1.00     1.00      75

   accuracy                           1.00      120
  macro avg       1.00     1.00     1.00      120
weighted avg       1.00     1.00     1.00      120

AUC Score: 1.00
```



5. K – Nearest Neighbors:

```
Best K-Nearest Params: {'metric': 'manhattan', 'n_neighbors': 11, 'weights': 'distance'}
```

```
===== K-Nearest Neighbors =====
```

```
Accuracy: 0.78
```

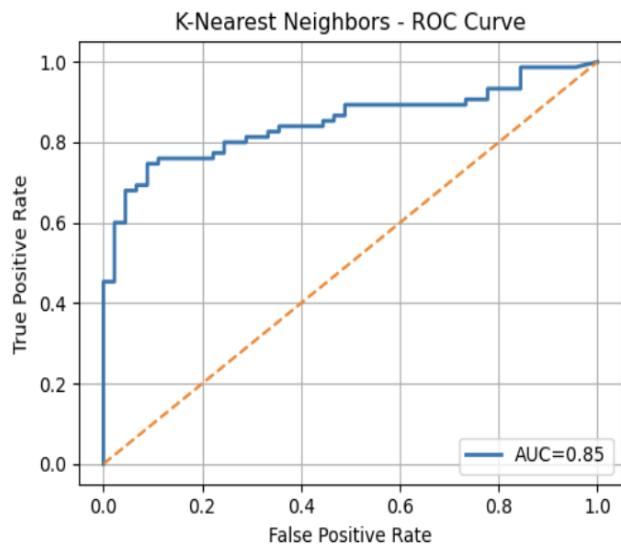
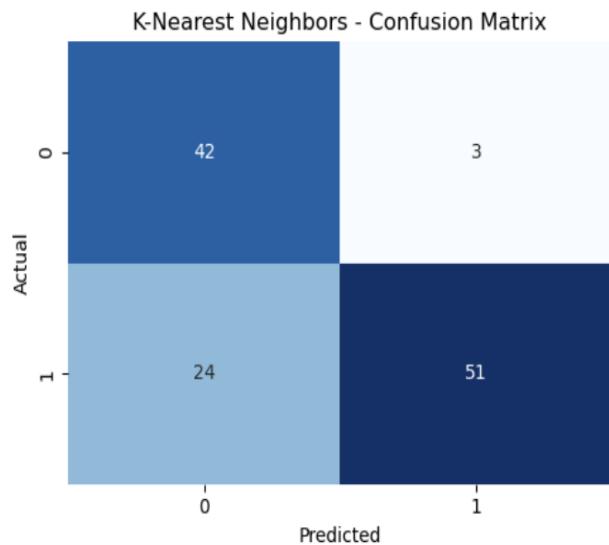
```
Confusion Matrix:
```

```
[[42  3]
 [24 51]]
```

```
Classification report:
```

	precision	recall	f1-score	support
0	0.64	0.93	0.76	45
1	0.94	0.68	0.79	75
accuracy			0.78	120
macro avg	0.79	0.81	0.77	120
weighted avg	0.83	0.78	0.78	120

```
AUC Score: 0.85
```

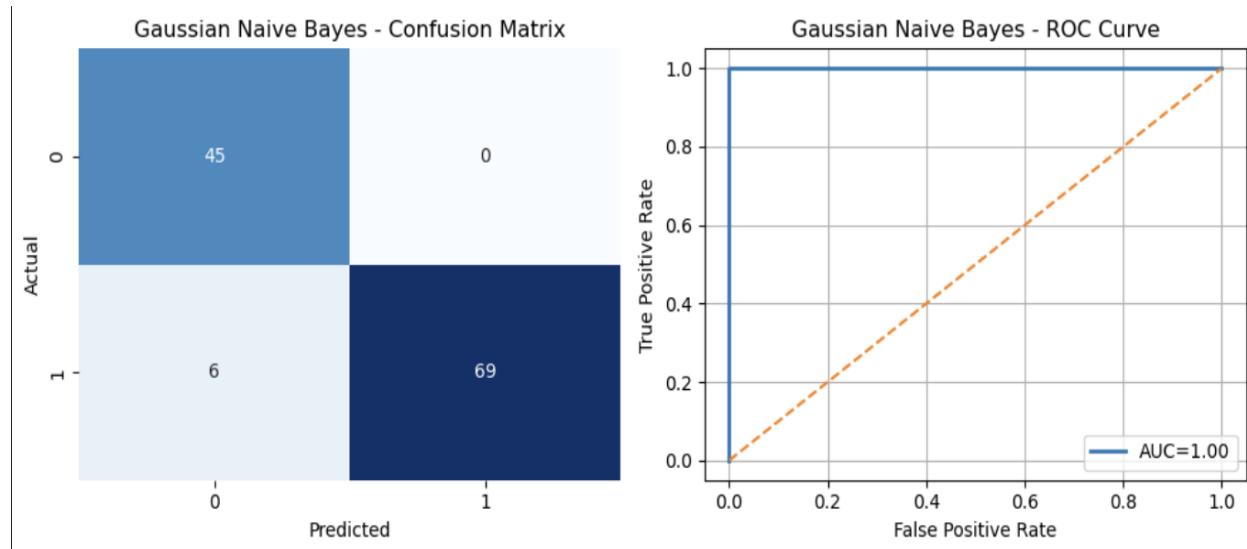


6. Gaussian Naïve Bayes:

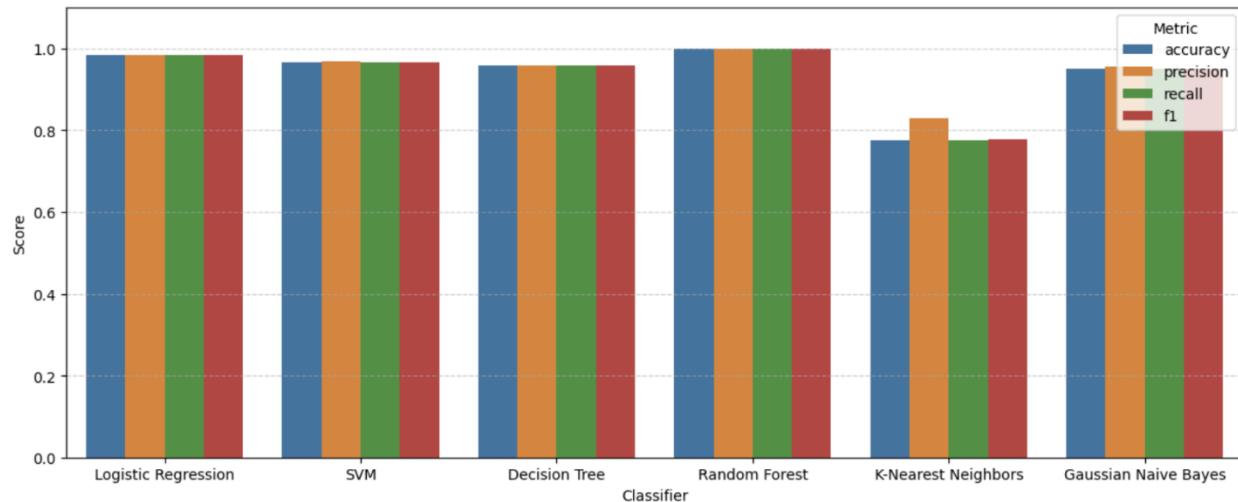
```
===== Gaussian Naive Bayes =====
Accuracy: 0.95
Confusion Matrix:
[[45  0]
 [ 6 69]]
Classification report:
precision    recall    f1-score   support
          0       0.88      1.00      0.94      45
          1       1.00      0.92      0.96      75

accuracy                           0.95      120
macro avg       0.94      0.96      0.95      120
weighted avg    0.96      0.95      0.95      120

AUC Score: 1.00
```



Comparison across evaluation metrics:



Final Model Selection:

Chosen Model: Random Forest Classifier

Justification:

- After experimenting with multiple classification algorithms, the **Random Forest model** achieved the best performance on the test set, with the highest **accuracy, F1-score, and ROC-AUC**, and excellent **recall** for the CKD class.
- Considering the clinical importance of correctly identifying CKD patients and the robustness of tree-based ensembles, the **Random Forest classifier** was chosen as the final model for this problem.

Summary of Model Performance:

Model	Accuracy	Precision	Recall	F1 score
Logistic Regression	0.98	0.98	0.98	0.98
SVM	0.97	0.97	0.97	0.97
Decision Tree	0.96	0.96	0.96	0.96
Random Forest	1.00	1.00	1.00	1.00
K- Nearest Neighbors	0.78	0.83	0.78	0.78
Naïve Bayes	0.95	0.96	0.95	0.95

